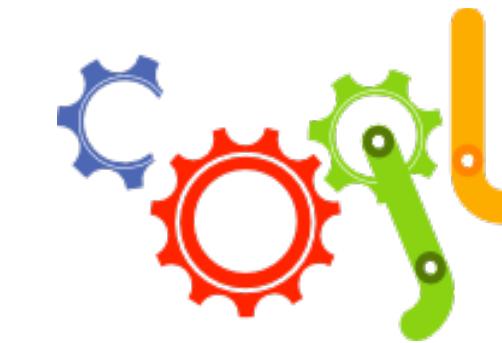
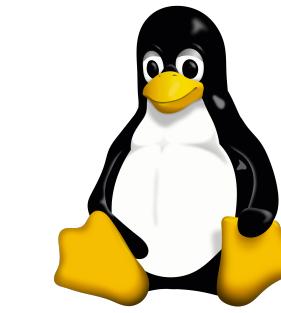


**Building an AI  
that Codes**

<http://chriscummins.cc>

2013



2014



2015



2016







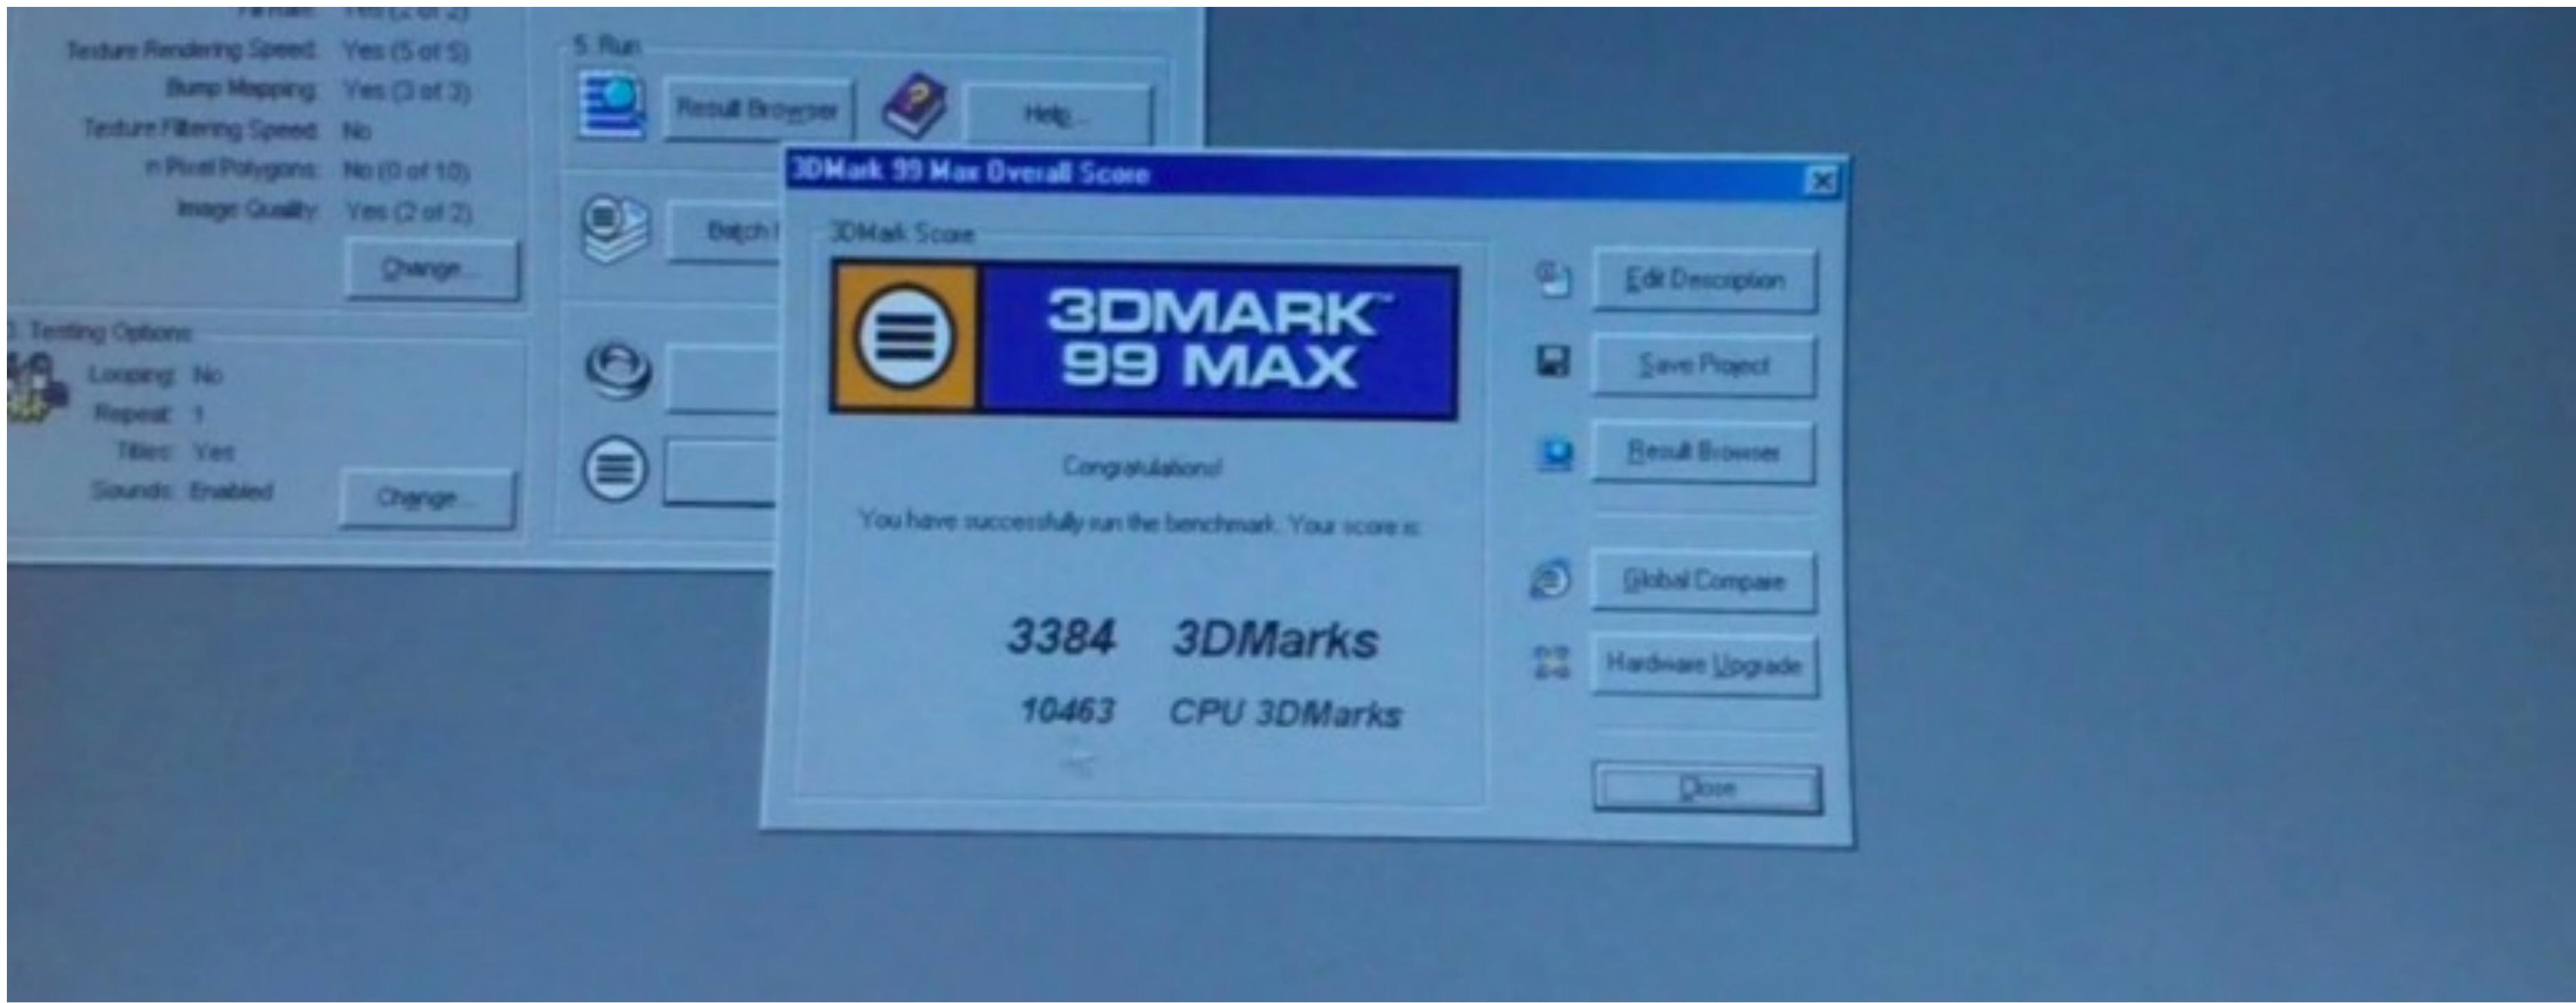
GIFTS

TECH-SPEC  
LACES

MOBILE ARREST  
CENTER

POLICE  
OF CHICAGO

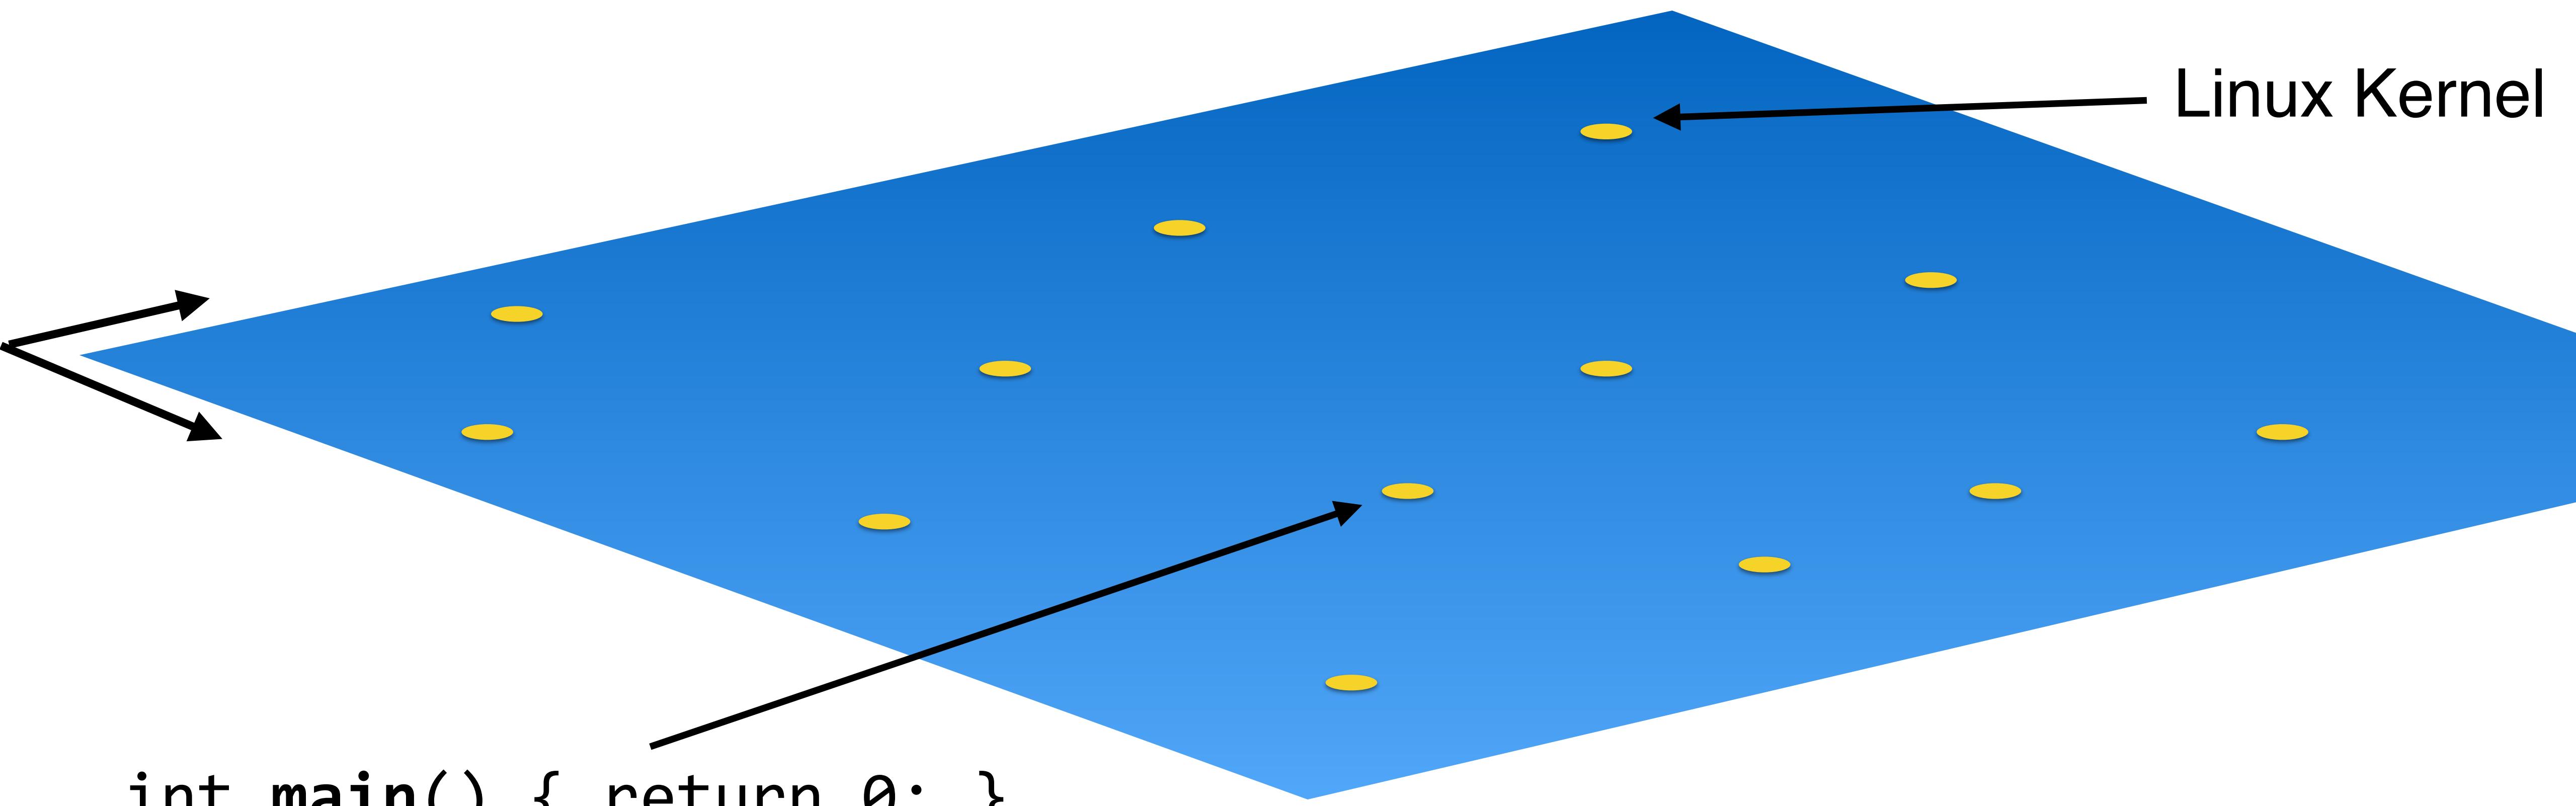
# What makes a good computer?



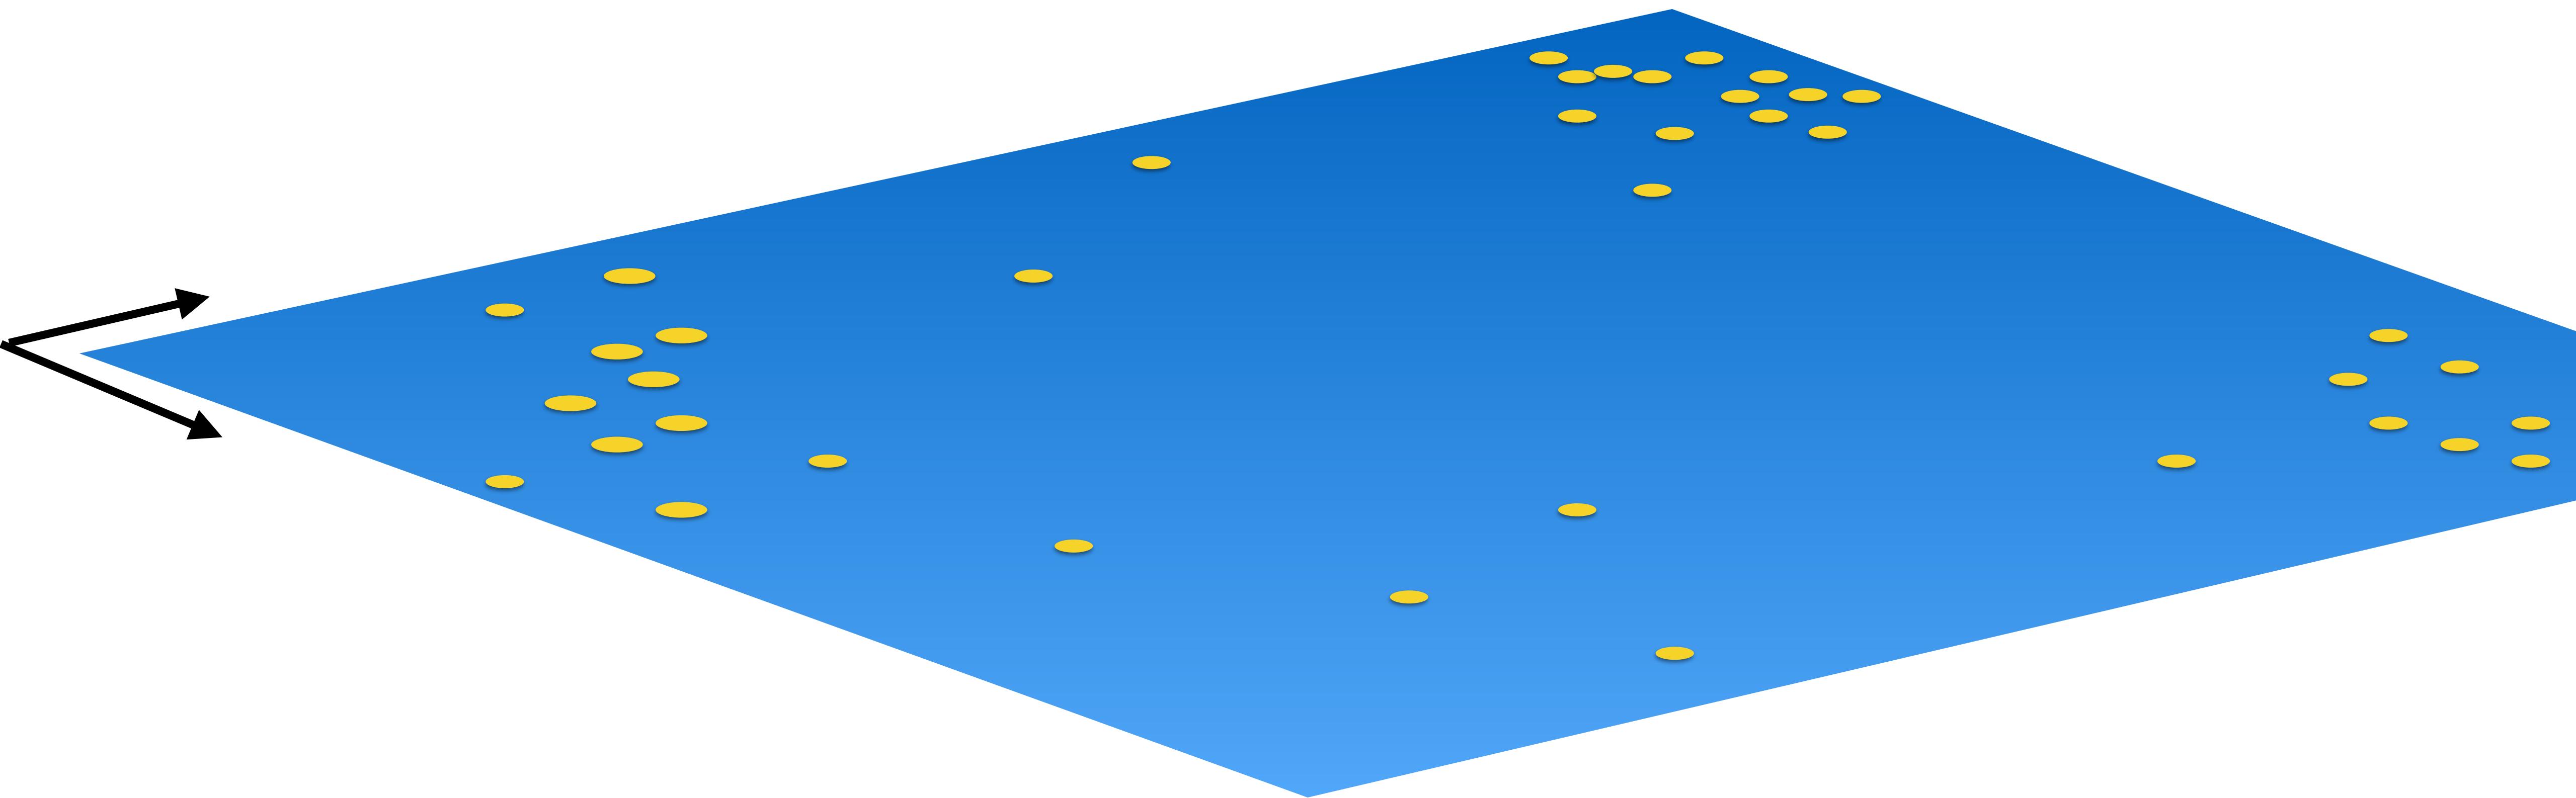
big numbers != smooth games

fast forward . . .

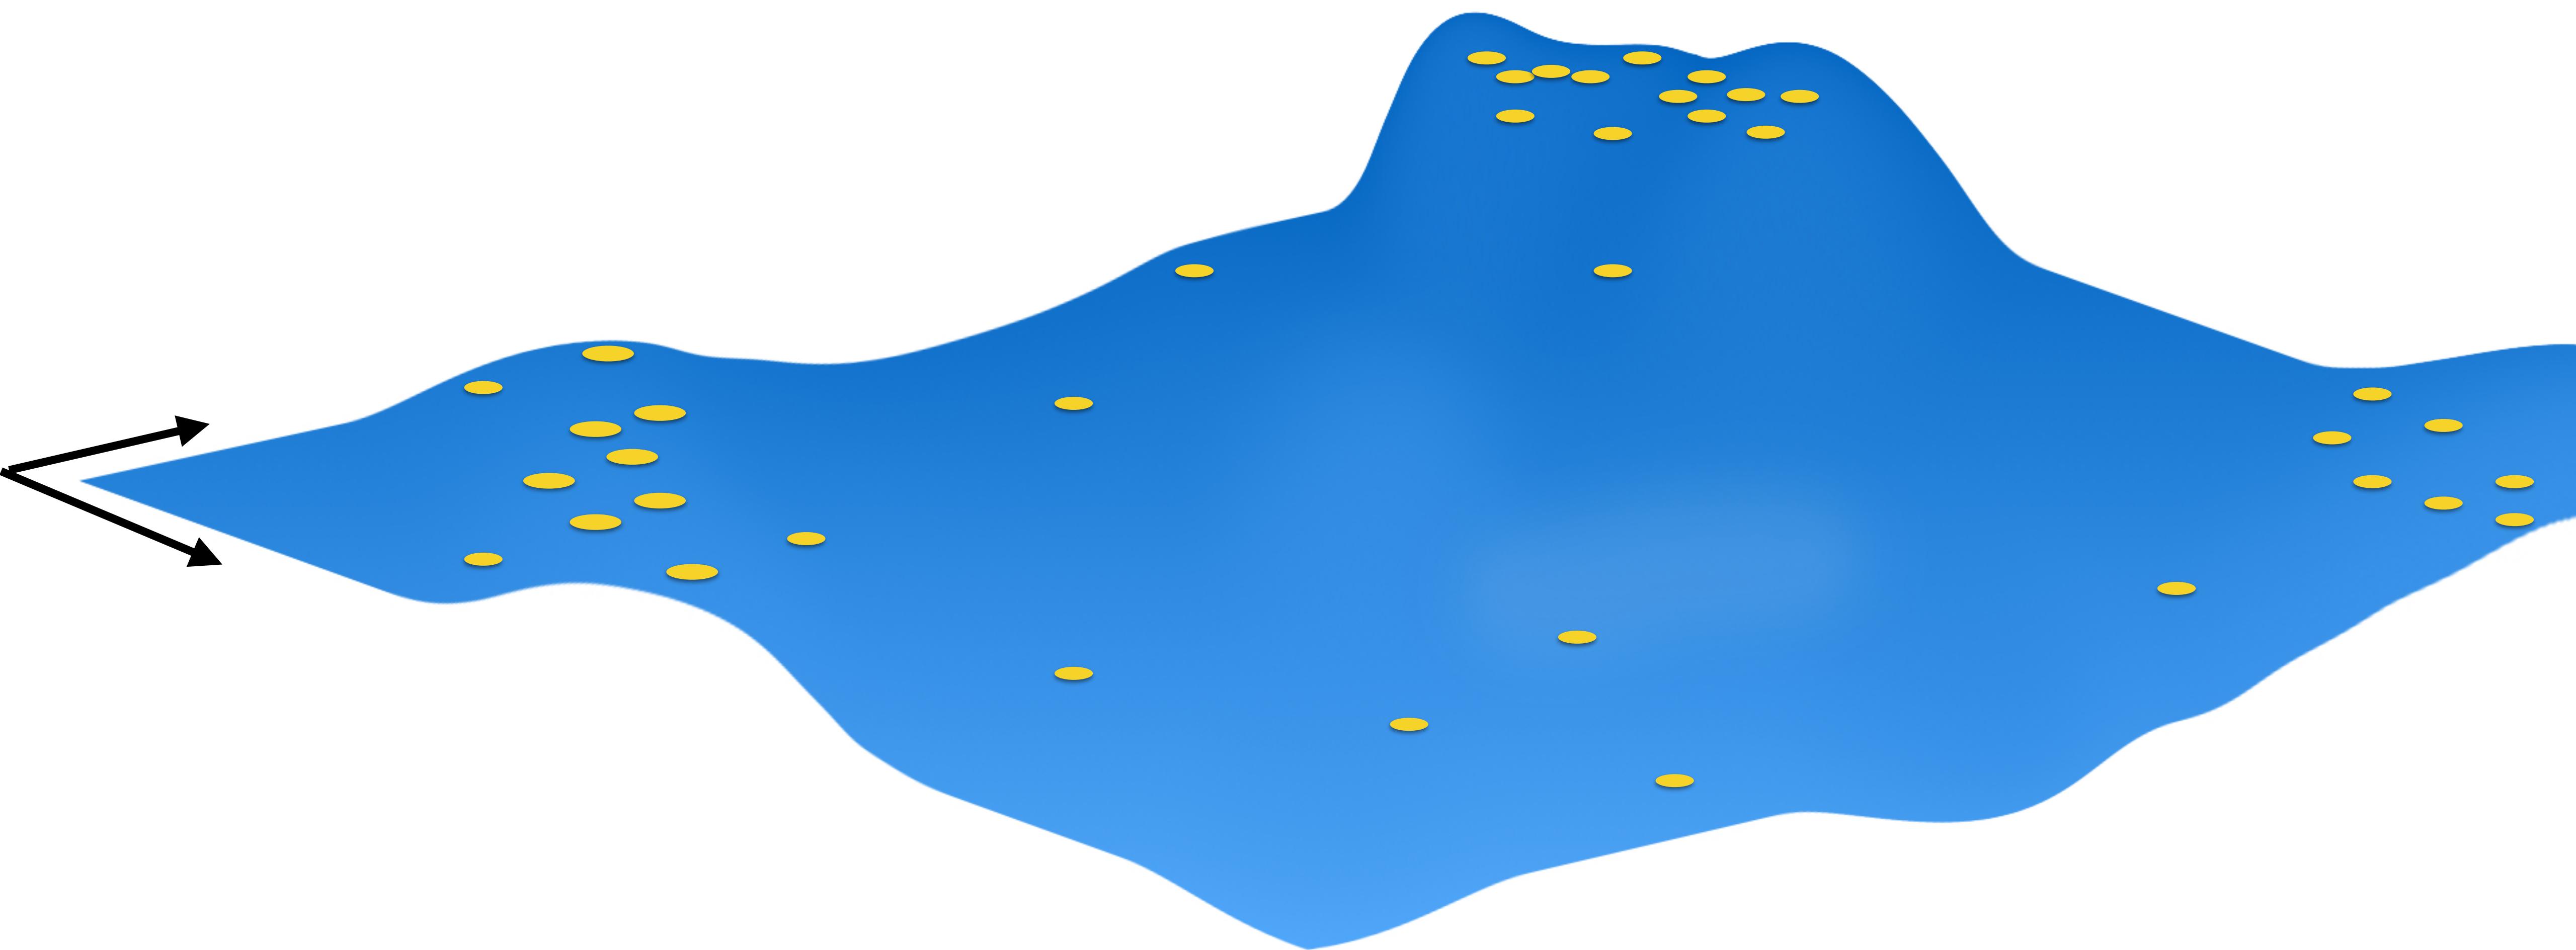
# Consider the “implementation space”



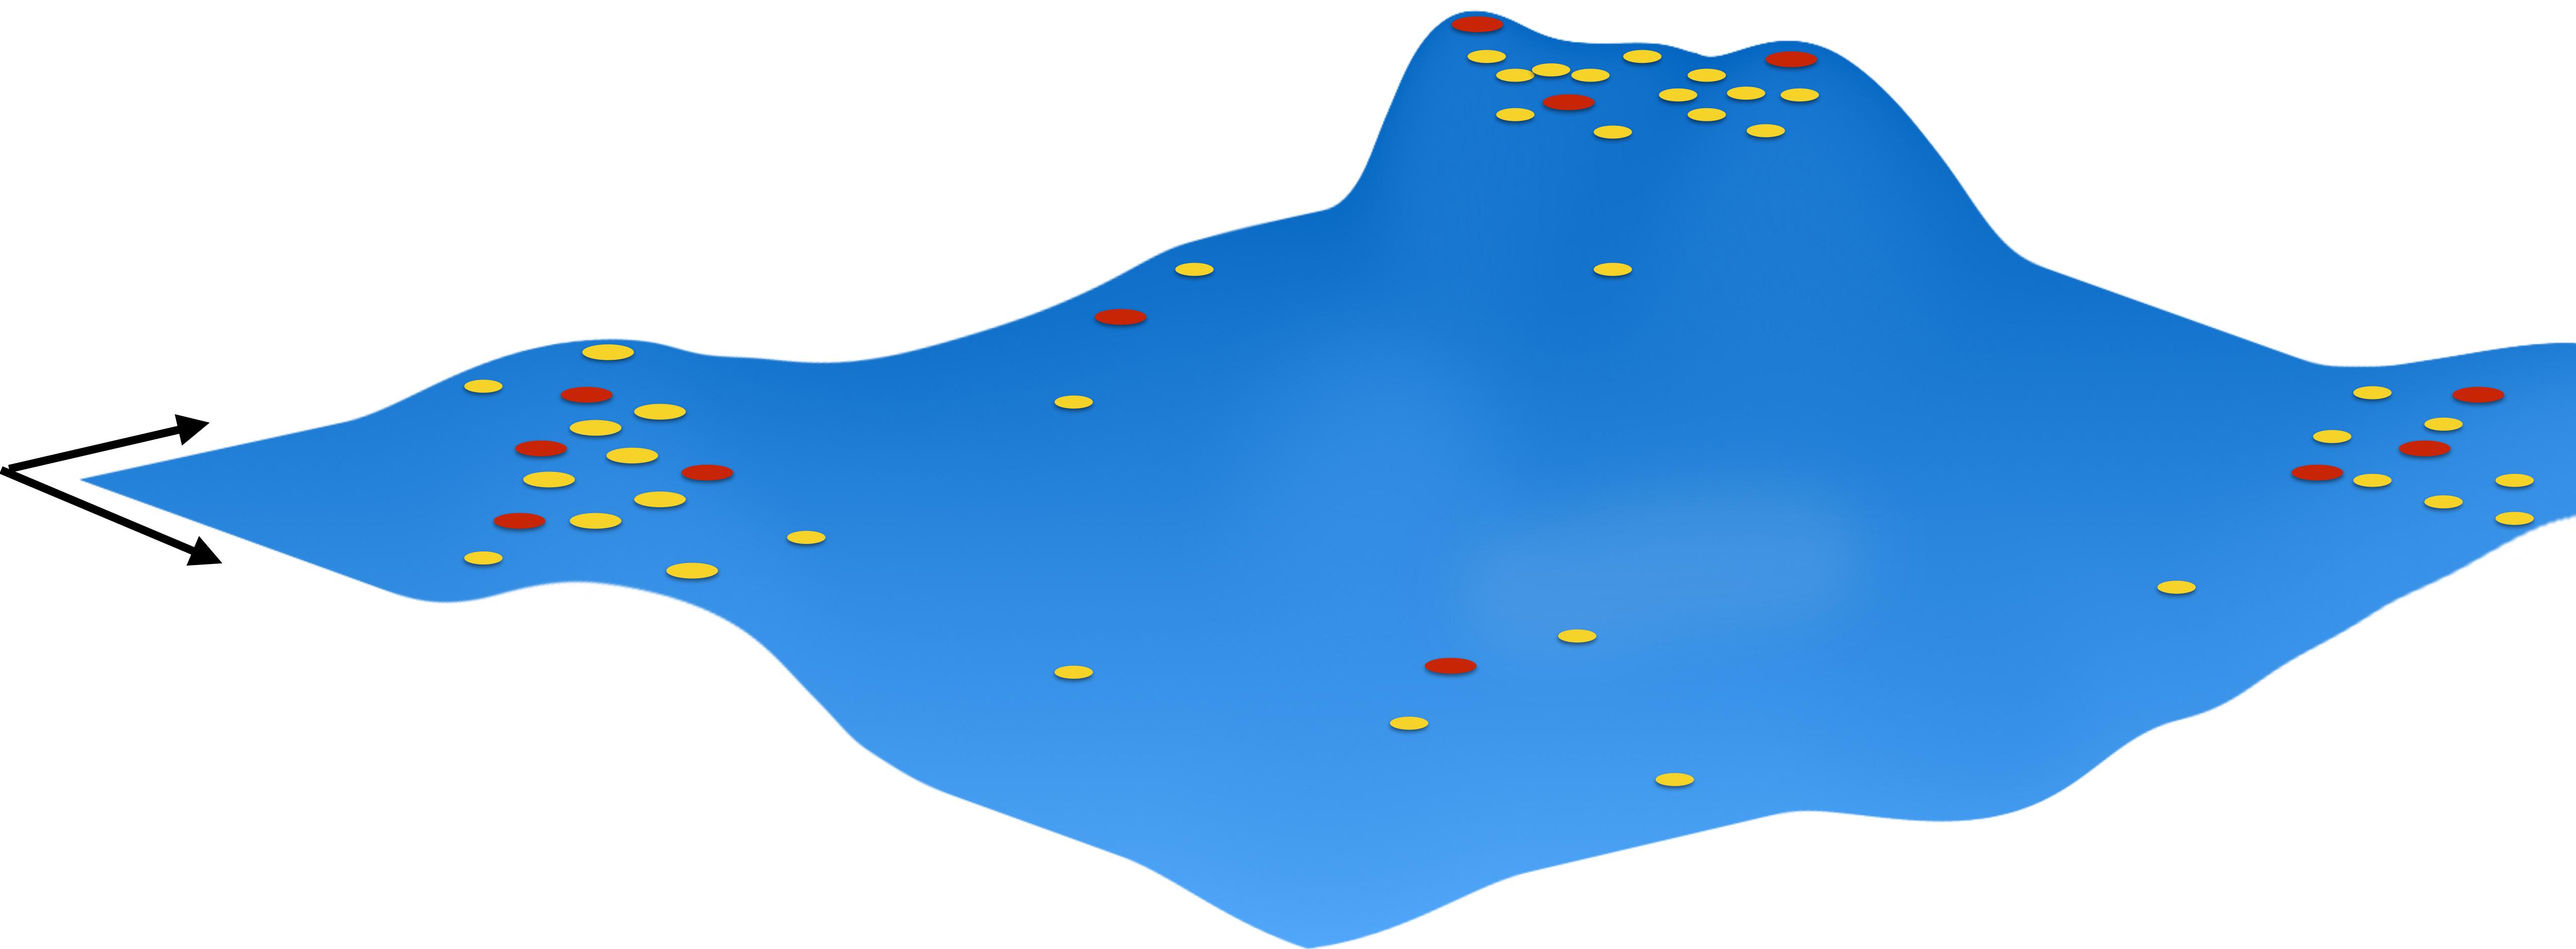
# Hypothesis: real source codes form clusters



# Weight space to match clustering



Sample from weighted space to generate new,  
representative benchmarks on-demand.



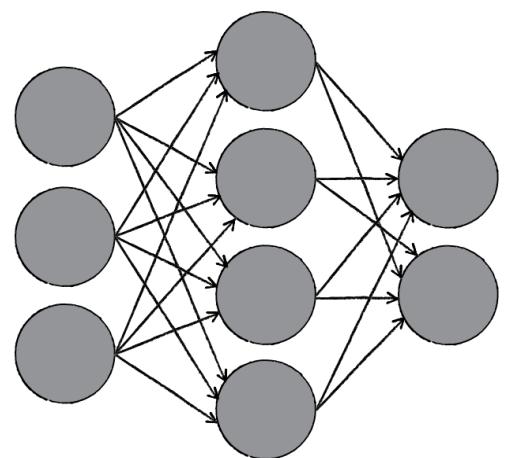
# The idea



Mine programs from



Apply



over implementation space

Generate representative benchmarks on-demand

Human or Robot? Chris

humanorrobot.uk/game/?g=opencl&m=rabt#

## humanorrobot.uk

**Round 1**

**Player: 1000, Robot: 1000**

```
_kernel void A(__global float *a, __global float *b, __global int *c) {
    int d = get_global_id(0);
    float16 e = (float16)db* (*c), db* (*c], db* (*c],
    db* (*c], db* (*c],
    db* (*c], db* (*c],
    db* (*c0], db* (*c1],
    db* (*c2], db* (*c3],
    db* (*c2], db* (*c3],
    db* (*c4], db* (*c5];
}

float16 f;
f = cosh(e);
a[d * (*c) + 0] = f[0];
a[d * (*c) + 1] = f[1];
a[d * (*c) + 2] = f[2];
a[d * (*c) + 3] = f[3];
}
```

```
_kernel void A(__global float *a, __global float *b, __global float *c) {
    float d;
    float e = 0;
    for (int f = 0; f < 1024; f++) e[e * 16 + g] = i[e * (*c) + 0];
    barrier(1);

    for (int g = 1; g < c * (g); ++g) {
        for (int h = 0; h < i - 1; h++) {
            g[h] = f[g] * ((h + 1 & 0xf) << (1 - i) | (i[h] >> (32 - h)));
        }
        h = (h & 0x80F) + (f - h) >> 1;
        i = f + (e + f) / 2;
        b[f] = i;
    }

    return (f & g);
}
```

**This is more human-like**

**This is more human-like**

© 2016 Chris Cummins.

# http://humanorrobot.uk

**Curing the Benchmark Deficit: On-Demand Compute Kernel Synthesis using Deep Learning**

Chris Cummins  
University of Edinburgh  
c.cummins@ed.ac.uk

Pavlos Petoumenos  
University of Edinburgh  
ppetoume@inf.ed.ac.uk

Zheng Wang  
Lancaster University  
z.wang@lancaster.ac.uk

Hugh Leather  
University of Edinburgh  
hleather@inf.ed.ac.uk

**Abstract**  
The quality of performance tuning is bound by the quantity and quality of benchmarks used. Too few benchmarks leads to overfitting; non-representative benchmarks lead to invalid predictions.  
We present a novel methodology for generating OpenCL compute kernels. Given a corpus of example programs, we apply deep learning across the implementation space, learning a language model from which we obtain new kernels through a process of rejection sampling. We demonstrate our approach for a state-of-the-art machine learning OpenCL autotuner. With the addition of synthesised compute kernels, we improve the accuracy of machine learning predictions from XXX% to XXX%, demonstrating up to X $\times$  speedup over the hand-selected benchmarks.

**Keywords** Synthetic program generation, OpenCL, Deep Learning, GPUs

**1. Introduction**  
Benchmarking parallel applications is hard. State of the practice is lacking [1]. OpenCL benchmark suites: Rodinia [2], Parboil [3], Polybench [4], SHOC [5], AMD SDK <sup>1</sup> and NVIDIA SDK <sup>2</sup>. TODO: Tease the small number of benchmarks used.

Benchmarking OpenCL [6]. Are benchmarks suites representative? Exploring the full performance spectrum [7]. Characterising workloads of Rodinia and Parsec [8].

In previous works we used stochastic template substitution [9, 10]. This template based approach is not general

<sup>1</sup><http://developer.amd.com/>  
<sup>2</sup><https://developer.nvidia.com/opencl/>

**RQ1:** Can the quality of machine learning predictions be improved with the addition of representative benchmarks?  
For this, we need a methodology for generating novel source codes. This leads to the subsequent research question:  
**RQ2:** Given a program checker and a corpus of example programs, can a language model learn to generate new programs?

The key contributions of this work are as follows:

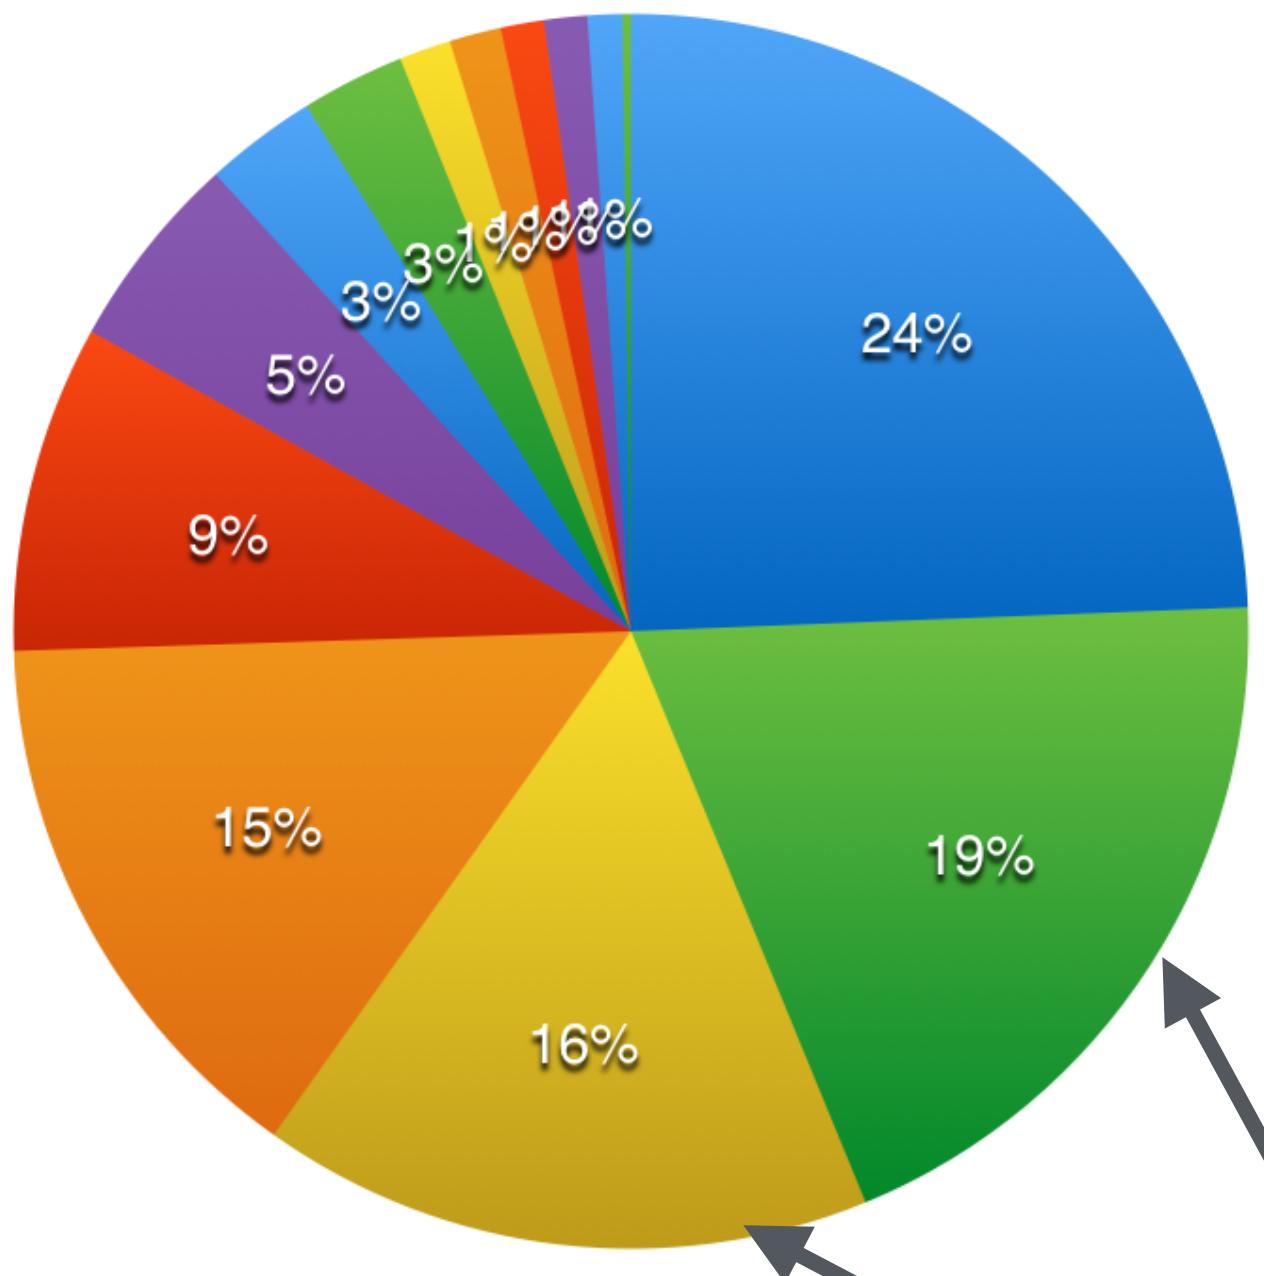
- a novel methodology for generating synthetic compute kernels for machine learning-based performance tuning;
- a large scale source code evaluation and re-writing for character-level language modelling using deep learning;
- improved tuning performance of [CGO'13] on hand written OpenCL kernels.

**2. Motivation**  
We surveyed the benchmarking methodologies of GPU research papers from top tier conferences between 2013–2016: CGO, HiPC, PACT, and PPoPP. By aggregating the sources and quantities of benchmark kernels from 27 papers, we find that 79% of benchmark kernels come from four benchmark suites <sup>3</sup>. Figure 1 and Figure 2.

**3. Generating Compute Kernels**  
Overview of methodology. Figure 3.

<sup>3</sup>Raw data available at: <http://bit.ly/foo>

2013-2016 state of practise  
27 top-tier GPU papers



# (SDK sample codes)